

# CS32 Intro to Computer Science II

---

Baoxiong Jia & Muthu Palaniappan, DIS 1C Week 9  
UCLA Spring 2021

# About Us

- TA: Baoxiong Jia
  - Email: [baoxiongjia@cs.ucla.edu](mailto:baoxiongjia@cs.ucla.edu)
  - Office Hours: Tuesday 8:30-10:30am
  - Thursday 8:30-10:30am
  - Discussion 1C: Friday 12:00-13:50pm
- LA: Muthu Palaniappan
  - Email: [muthupal@g.ucla.edu](mailto:muthupal@g.ucla.edu)
  - Office Hours: Monday 10:30-11:30am
  - Wednesday 10:30-11:30am

# Outline

- Hash tables

# Hash tables

- Hashing
  - Distribute the entries (key and value pairs) across an array of buckets by using a hash function
- Hash function
  - A function carefully chosen to provide a uniform distribution of hash values to generate keys
  - The key is to avoid collisions of keys generated

# Hashing

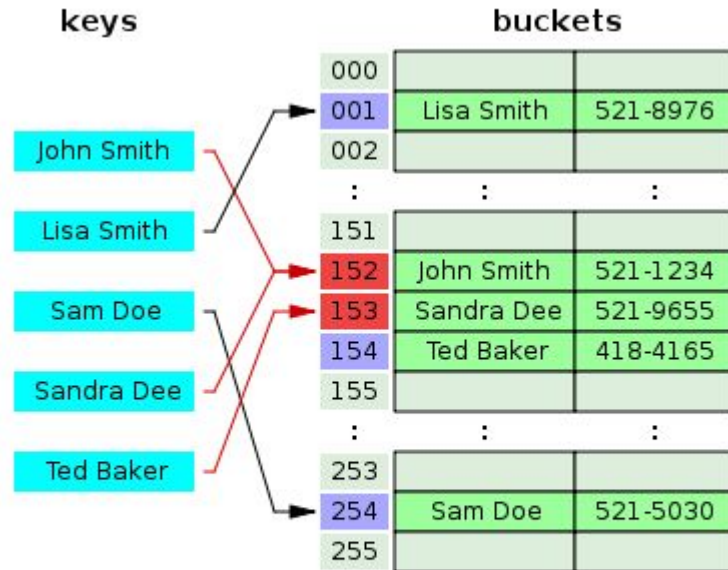
- Collision could occur even when the number of buckets is significantly larger than the number of entries
  - 10000 buckets to hash a dataset of strings, e.g. {Amy, Brendon, Cindy, David, ...}
  - A hash function  $f(x)$  based on capital letters where  $f(\text{'Amy'}) = \text{'a'}$ ,  $f(\text{'Annie'}) = \text{'a'}$ , ...
- Load factor

$$\text{load factor} = \frac{n}{k},$$

where

- $n$  is the number of entries occupied in the hash table.
- $k$  is the number of buckets.

# Resolving collision



Hash collision resolved by open addressing with linear probing (interval=1). Note that "Ted Baker" has a unique hash, but nevertheless collided with "Sandra Dee", that had previously collided with "John Smith".

To resolve collision, we typically use:

For a closed hash table

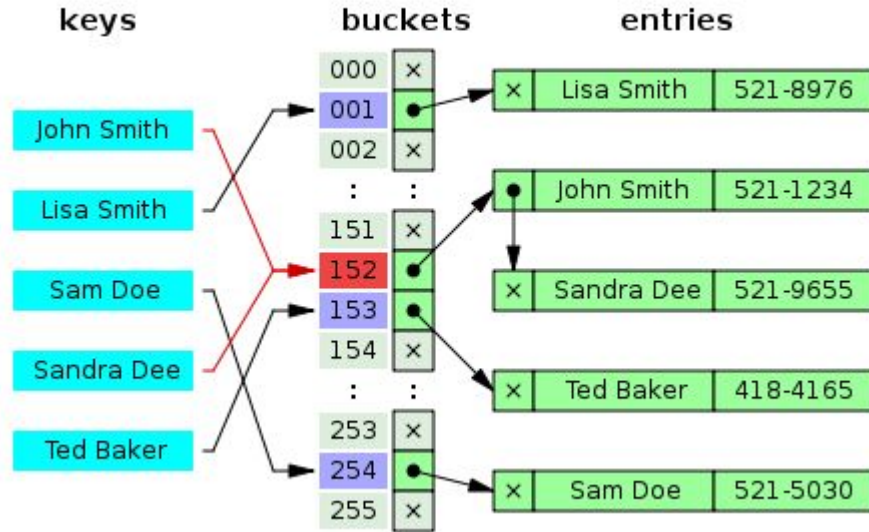
- Open addressing
  - Linear probing

Our **closed hash table** + **linear probing** works just fine, but it still has a few problems:

It's **difficult** to **delete** items

It has a **cap** on the **number of items** it can hold... **That's a bummer.**

# Resolving collision



To resolve collision, we typically use:

For an open hash table

- Separate chaining
  - Linked list
  - Binary search trees

# Hash table in C++

<https://repl.it/@jiajerry/HashTable#main.cpp>



class template

**std::unordered\_map**

<unordered\_map>

```
template < class Key,                                // unordered_map::key_type
           class T,                                  // unordered_map::mapped_type
           class Hash = hash<Key>,                   // unordered_map::hasher
           class Pred = equal_to<Key>,                // unordered_map::key_equal
           class Alloc = allocator< pair<const Key,T> > // unordered_map::allocator_type
       > class unordered_map;
```

## Unordered Map

Unordered maps are associative containers that store elements formed by the combination of a *key value* and a *mapped value*, and which allows for fast retrieval of individual elements based on their keys.

In an `unordered_map`, the *key value* is generally used to uniquely identify the element, while the *mapped value* is an object with the content associated to this key. Types of *key* and *mapped value* may differ.

Internally, the elements in the `unordered_map` are not sorted in any particular order with respect to either their *key* or *mapped* values, but organized into *buckets* depending on their hash values to allow for fast access to individual elements directly by their *key values* (with a constant average time complexity on average).

`unordered_map` containers are faster than `map` containers to access individual elements by their *key*, although they are generally less efficient for range iteration through a subset of their elements.

Unordered maps implement the direct access operator (`operator[]`) which allows for direct access of the *mapped value* using its *key value* as argument.

Iterators in the container are at least *forward* iterators.